

Greedy Technique


- Constructs a solution to an *optimization problem* piece by piece through a sequence of choices that are:
 - *feasible*
 - *locally optimal*
 - *irrevocable*
- For some problems, yields an optimal solution for every instance.
- For some problems where it does not yield an optimal solution it can be useful for fast approximations

Applications of the Greedy Strategy

- Optimal solutions:
 - change making for “normal” coin denominations
 - minimum spanning tree (MST)
 - single-source shortest paths
 - simple scheduling problems
 - Huffman codes
- Approximations:
 - traveling salesman problem (TSP)
 - knapsack problem
 - other combinatorial optimization problems

Interval Scheduling

- Interval scheduling.
 - Job j starts at s_j and finishes at f_j .
 - Two jobs compatible if they don't overlap.
 - Goal: find maximum subset of mutually compatible jobs.



Interval Scheduling: Greedy Algorithms

- Greedy template. Consider jobs in some order. Take each job provided it's compatible with the ones already taken.
 - [Earliest start time] Consider jobs in ascending order of start time s_j .
 - [Earliest finish time] Consider jobs in ascending order of finish time f_j .
 - [Shortest interval] Consider jobs in ascending order of interval length $f_j - s_j$.
 - [Fewest conflicts] For each job, count the number of conflicting jobs c_j . Schedule in ascending order of conflicts c_j .

Interval Scheduling: Greedy Algorithms

- Greedy template. Consider jobs in some order. Take each job provided it's compatible with the ones already taken.



breaks earliest start time



breaks shortest interval



breaks fewest conflicts

Interval Scheduling: Greedy Algorithm

- Greedy algorithm. Consider jobs in increasing order of finish time. Take each job provided it's compatible with the ones already taken. (Earliest finish time first)

```
Sort jobs by finish times so that  $f_1 \leq f_2 \leq \dots \leq f_n$ .
jobs selected
A ← ∅
for j = 1 to n {
    if (job j compatible with A)
        A ← A ∪ {j}
}
return A
```

- Implementation. $O(n \log n)$.
 - Remember job j^* that was added last to A.
 - Job j is compatible with A if $s_j \geq f_{j^*}$.


Interval Scheduling – Stay Ahead proofs

Proof of correctness

- See handout on PolyLearn for a proof by induction


Interval Scheduling: Theorem. Greedy (E.F.T.F) algorithm is optimal

- Pf. (by contradiction)
 - Assume greedy is not optimal, and let's see what happens.
 - Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.
 - Let j_1, j_2, \dots, j_m denote set of jobs in an optimal solution which has the largest set initial set of jobs that are the same as those in the greedy solution, $m > k$. The figure below shows this assuming the first $r < k$ jobs are the same.
 - Since the i_{r+1} job has the earliest finish time compatible with $i_r = j_r$ it must finish before j_{r+1} . Replace j_{r+1} with i_{r+1} and the remaining jobs in the optimal solution will still be compatible with $i_1 \dots i_{r+1}$.



Theorem. Greedy algorithm is optimal.

- Pf. (continued)
 - This forms a solution of m compatible jobs that agree with the greedy solution for the first $r+1$ jobs .
 - This contradicts the assumption that the longest initial agreement between an optimal solution with the greedy was $r < k$ jobs.
 - Since there is an optimal solution with $r=k$ we have the situation below. But it is impossible that there are remaining jobs compatible with job k since greedy would have chose one.



Total Weighted Time to Completion Problem

Total Weighted Time to Completion Problem: Given a set of jobs: $job_1, job_2, \dots, job_n$ characterized by their lengths: l_1, l_2, \dots, l_n and their weights: w_1, w_2, \dots, w_n

- Output a schedule that minimizes the total weighted time to completion.
 - Schedule: An ordering of the jobs to be serviced by the resource
 - Completion time: C_i , of job i is the sum of the l_j 's up to and including i
- Total Weighted Time to Completion = $\sum_{i=1 \text{ to } n} w_i C_i$

Question?

- Given 3 jobs job_1 , job_2 , job_3 where the respective lengths are 1, 2, 3. What are the Completion times of the three jobs if done in the order given?
 - a) 1, 2, 3
 - b) 3, 5, 6
 - c) 1, 3, 6
 - d) 1, 4, 6
- If the weights are respectively 3, 2, 1. What is the Total Weighted Time to Completion for the ordering 1, 2, 3 ?

Note: If there are n jobs then there are $n!$ orderings.

Total Weighted Time to Completion Problem (TWTC)

Finding a Greedy algorithm that solves this problem:

Focus on process

- Goal: Find a sequence of a set of jobs that **minimizes** the total weighted time to completion.
- There may be many candidate greedy algorithms!
- How do we find a greedy algorithm that works.
 - Look at special cases
 - Generate candidate greedy algorithms
 - Narrow down to single (if possible) candidate, usually by looking at simple cases to eliminate candidates, then prove correctness. (It may be that no greedy algorithm works.)

Discovery of candidates

- Explore special cases:
 - Equal lengths
 - Equal weights
- What order do you think you should schedule jobs?
 - a) Larger weights first; shorter lengths first
 - b) Smaller weights first; shorter lengths first
 - c) Larger weights first; longer lengths first
 - d) Smaller weights first; longer lengths first
- Hint: Look at small examples to test your ideas

What are reasonable candidates?

- How can we combine the information we have into a greedy algorithm
- Two simple options, order the jobs by either:
 - Difference: $w_i - l_i$ Ratio: w_i / l_i
- How to rule one of these out?
 - Find, if possible, an example where the two different algorithms give two different answers.
 - Clearly the larger solution for the example is not optimal.
 - This rules out that algorithm but does not prove anything about the algorithm that gives the smaller solution.

TWTC example problem

Consider 2 jobs $l_1 = 1, w_1 = 5$ $l_2 = 4, w_2 = 12$

- What is the weighted total time to completion given by the two algorithms?

2 jobs $l_1 = 1, w_1 = 5$ $l_2 = 4, w_2 = 12$

- What is the weighted total time to completion given by the two algorithms?

	$W_i - L_i$	W_i/L_i
Job ₁	4	5
Job ₂	8	3
Ordering	Job ₂ , Job ₁	Job ₁ , Job ₂
TWTC	$48+25 = 73$	$5+60 = 65$

- $W_i - L_i$ does not always yield the optimal (smallest) TWTC.

Proving correctness: exchange argument

Plan: Assume the greedy algorithm does not produce the optimal solution.

- Assume the w/l ratios are all distinct so there is only one possible ordering given by the algorithm.
- Assume that the jobs are labeled according to the decreasing w/l ratio, e.g. w_1/l_1 is the largest ratio, w_n/l_n is smallest
- Assume that there is an ordering Σ^* that is the optimal solution but is different from Σ the ordering given by the algorithm

Now show that Σ^* is not the optimal solution.

This contradicts our assumption that Σ^* is optimal and different from Σ . Hence the optimal solution must be Σ .

Proving correctness: exchange argument details

the jobs are labeled according to the decreasing ratios

If Σ^* is not Σ , it must have two consecutive jobs where $i > j$ (i has the smaller ratio) but i comes before j in Σ^* .

Why must there be two consecutive jobs out of order ?

Order of jobs in Σ^*

Jobs before i	Job i	Job j	Jobs after j
-----------------	---------	---------	----------------

If switch jobs i and j , what happens to the completion times of jobs?

Proving correctness: exchange argument details

Jobs before i	Job i	Job j	Jobs after j
---------------	-------	-------	--------------

If switch jobs i and j, what happens to the completion times of jobs?

$k \neq i, j$	i	j
– unaffected	up	down

- Now what does switching i and j do to the total cost?
 - Since $i > j$ we know that $w_i / l_i < w_j / l_j$ so
 - $w_i * l_j < w_j * l_i$ means that $w_j * l_i - w_i * l_j > 0$
 - Claim: the ordering after making the switch has a lower total weighted time to completion! CONTRADICTION that Σ^* was the minimum total weighted time to completion

details

Before	Jobs before i	Job i	Job j	Jobs after j
	C			

After	Jobs before i	Job j	Job i	Jobs after j
	C			

- Weighted Completion time:

After	job _i	$w_i * (C + l_j + l_i)$	job _j	$w_j * (C + l_j)$
Before	job _i	$w_i * (C + l_i)$	job _j	$w_j * (C + l_i + l_j)$

Before – After = $w_j * l_i - w_i * l_j$ but this is > 0


\Rightarrow After has a smaller TWTC

Contradicts “Before” being optimal \Rightarrow ordering by decreasing weight to length ratio gives an optimal ordering


Minimum Spanning Tree Problem

- Spanning tree of a connected graph G : a connected acyclic subgraph of G that includes all of G 's vertices
- Minimum spanning tree, MST , of a weighted, connected graph G : a spanning tree of G of minimum total weight (sum of the edge weights in the tree)
- Minimum Spanning Tree Problem: For a weighted graph, find the MST

Minimal Spanning Tree Example



Minimal Spanning Tree Example



3 Greedy Approaches – these all work!

- Build a spanning tree by
 - Successively adding the lightest edge that does not create a cycle (Kruskal's)
 - Start with a root node and grow the tree outward without creating a cycle (Prim's)
 - Start with the full graph and delete edges in order of decreasing cost as long as we do not disconnect the tree (Reverse-delete?)

Prim's MST algorithm

- Start with tree T_1 consisting of one (any) vertex and “grow” tree one vertex at a time to produce MST through a series of expanding subtrees T_1, T_2, \dots, T_n
- On each iteration, construct T_{i+1} from T_i by adding vertex not in T_i that is closest to those already in T_i (this is a “greedy” step!)
 - Key Question – how do we do this efficiently
- Stop when all vertices are included

Cut property – proving MST algorithms correct

Assume all the edge costs are distinct. Let

- S be non-empty proper subset of V
- $e = (v, w)$ the minimal cost edge between S and $V-S$
- then every MST contains e




FIGURE 9.4 Correctness proof of Prim's algorithm.

Cut property – proving MST algorithms correct

Assume all the edge costs are distinct. Let

- S be non-empty proper subset of V
- $e = (v, w)$ the minimal cost edge between S and $V-S$
- then every MST contains e

Proof: (sketch – fill in the details)

- Let T be a spanning tree that does not contain e
- Use an exchange argument
- Find an e' so that exchanging e for e' reduces the cost
 - » There must be a path from v to w in T
 - » Follow the path to find an edge that goes from S to $S-V$
 - » Let that be e' , exchange to get T'
 - » Finally show: T' is a tree (acyclic) and has lower cost

Correctness: Prim's Algorithm


Correctness : Use induction to show that all the T_i are subtrees of the MST (apply the cut property)

- Base case: $T_1 = \{\text{start vertex}\}$. Clearly a subtree of a Minimal Spanning of T
- Inductive case: T_k a subtree of an MST $\rightarrow T_{k+1}$ a subtree of an MST
- Proof of inductive case: Apply the cut property to show that the edge chosen by Prim must be in any MST.
 - Let S be the set of vertices in T_k ; $V-T_k$ be will be the vertices not in T_k
 - By the cut property the smallest edge connecting T_k and $V-T_k$ must be in the MST
 - But this is the edge chosen by Prim's Algorithm to add to T_k to make T_{k+1}
 - This proves the inductive set – so all the subtrees created by Prim are subtrees of the MST
 - Hence all the edges in T_{k+1} are in the MST for the graph

Analysis: Prim's Algorithm

- Analysis: Using a heap implementation of a **priority queue** the algorithm will run:
 - $O(|E|)$ since each edge will eventually be placed in the queue
 - $|V| * \text{extractMin}$ and $|E| * \text{changePriority}$ operations
 - A priority queue can be implemented using a heap to so that both extractMin and changePriority are $O(\log |V|)$
 - Therefore the overall running time is $|E| * (\log |V|)$
- Why is it not proper to use the priority queue operations as the basic operations?

Example of Prim's Algorithm in action



Prim's Algorithm: More details

```

Initialize prev(v), dist(v)  $\forall v \in V$ ;  $T = \{\text{start vertex}\}$ ;  $U = V - T$ ; dist(start vertex) = 0
while(T is missing a vertex)
    pick the vertex, v1, in U with the shortest edge, dist(), to the group of vertices in
    the spanning tree add v1 to T
    /* this loop looks through every neighbor of v and checks to see if that
    * neighbor could reach the minimum spanning tree more cheaply through v1 */
    for each edge of v1 (v1, v2)
        if(length(v1, v2) < dist[v2]) // this is "relaxation"
            dist[v2] = length(v1, v2)
            prev[v2] = v1
        end if
    end for
end while

```

Single Source Shortest Paths Problem

Single Source Shortest Paths Problem: Given a weighted connected graph G , find shortest (sum of the weights of the edges) paths from source vertex s to each of the other vertices

- Variants:
 - Single destination shortest paths problem
 - Single pair shortest paths problem
 - All pairs shortest paths problem
- Issues
 - Negative edge weights
 - Cycles

Three Shortest Path Algorithms

Depend on idea of relaxation

- Dijkstra
 - Edge weights must be non-negative
 - Relaxes each edge exactly once. $O((\log |V|) * |E|)$
- Bellman-Ford
 - Works with negative edge weights (Of course not cycles!)
 - Relaxes edges multiple times $O(|V| * |E|)$
- DAG - assumes no cycles!
 - Uses topological sort of vertices
 - Guarantees $O(|E| + |V|)$ performance
 - Works with negative edge weights

Relaxation:

- For each vertex, v , maintain an upper bound on the length (stored in $\text{dist}(v)$) of the shortest path from the source vertex to that vertex and the previous vertex on that shortest path
- Test to see if some new path is shorter than the current upper bound. If so then reduce the upper bound and update the previous vertex.

Relax($e=(u,v)$, w)

if $\text{dist}(v) > \text{dist}(u) + w(u,v)$

$\text{dist}(v) = \text{dist}(u) + w(u,v)$

$\text{prev}(v) = u$


Dijkstra's Algorithm

■ Dijkstra's algorithm.

- Maintain a set of **explored nodes** S for which we have determined the shortest path distance $d(u)$ from s to u .
- Initialize $S = \{s\}$, $d(s) = 0$.
- Repeatedly choose unexplored node v which minimizes

$$\pi(v) = \min_{e = (u,v) : u \in S} d(u) + \ell_e, \quad \text{shortest path to some } u \text{ in explored part, followed by a single edge } (u, v)$$

add v to S , and set $d(v) = \pi(v)$.




Dijkstra's Algorithm

■ Dijkstra's algorithm.

- Maintain a set of **explored nodes** S for which we have determined the shortest path distance $d(u)$ from s to u .
- Initialize $S = \{s\}$, $d(s) = 0$.
- Repeatedly choose unexplored node v which minimizes

$$\pi(v) = \min_{e = (u,v) : u \in S} d(u) + \ell_e, \quad \text{shortest path to some } u \text{ in explored part, followed by a single edge } (u, v)$$

add v to S , and set $d(v) = \pi(v)$.



Representing Shortest Paths

- For each vertex u maintain the previous vertex from the source to that vertex, $\text{prev}(u)$
- Can reconstruct the shortest path to any vertex by walking backwards from the vertex to the source

```
Initialize_Single_Source(V)
  for  $v \in V$ 
     $\text{dist}(v) = \infty$ 
     $\text{prev}(v) = \text{null}$ 
   $\text{dist}(s) = 0$ 
```

Dijkstra's algorithm pseudo-code single source all shortest paths

```
For every vertex  $\text{dist}[v] \leftarrow \infty$  as distance;  $\text{prev}[v] = \text{null}$ ;
 $\text{Dist}[\text{source}] = 0$ ;
 $Q = V$  ( $Q$  will be the set of unfinished vertices)
 $V_T \neq \emptyset$ 
while  $Q$  is not empty.
  // Loop invariant:  $V_T = \{v: \text{know shortest path}\}$ 
  // each time through finds shortest path to vertex
  // and adds edge to  $V_T$ 
```

Dijkstra's algorithm pseudo-code

single source all shortest paths

- For every vertex $\text{dist}[v] \leftarrow \infty$ as distance; $\text{prev}[v] = \text{null}$;
- $\text{Dist}[\text{source}] = 0$;
- $Q = V$ (Q will be the set of unfinished vertices)
- $V_T \neq \emptyset$
- while Q is not empty. // Loop invariant: $V_T = \{v: \text{know shortest path}\}$
 // each time through finds shortest path to vertex and adds to V_T
 - u be vertex with smallest distance to source in Q , remove u^* from Q
 - $V_T = V_T \cup \{u^*\}$
 - for all of its unfinished neighbors v
 - » calculate their tentative distance to source through current node --- u^* .
 - » If this distance is less than the previously recorded tentative distance of v , then overwrite $\text{dist}[v]$ and update $\text{prev}[v]$

Algorithm: Prim-MST (G)

Input: Graph $G=(V,E)$ with edge-weights.

// Initialize priorities and place in priority queue.

1. $\text{priority}[i] = \text{infinity}$ for each vertex i

2. Insert vertices and priorities into priorityQueue ;

// Set the priority of vertex 0 to 0.

3. $\text{priorityQueue.decreaseKey}(0, 0)$ //

// Process vertices one by one in order of priority

```

4. while priorityQueue.notEmpty()
    // Get "best" vertex out of queue.
5.    $v = \text{priorityQueue.extractMin}()$ 
6.   Add  $v$  to MST;
    // Explore edges from  $v$ .
7.   for each edge  $e=(v, u)$  in  $\text{adjList}[v]$ 
8.      $w = \text{weight of edge } e=(v, u)$ ;
    // If it's shorter to get to MST via  $v$ , then update.
9.     if  $\text{priority}[u] > w$ 
10.       $\text{priorityQueue.decreaseKey}(u, w)$ 
11.       $\text{predecessor}[u] = v$ 
12.    endif
13.  endfor
14. endwhile

```

15. Build MST;

16. return MST

Algorithm: Dijkstra-SPT (G, s)

Input: Graph $G=(V,E)$ with edge weights ≥ 0

// Initialize priorities & put in priority queue.

1. $\text{priority}[i] = \text{infinity}$ for each vertex i ;

2. Insert vertices and priorities into priorityQueue ;

// Source s has priority 0

3. $\text{priorityQueue.decreaseKey}(s, 0)$

// Process vertices one by one in order of priority

```


4. while priorityQueue.notEmpty()
    // Get "best" vertex out of queue.
5.    $v = \text{priorityQueue.extractMin}()$ 
6.   Add  $v$  to SPT;
    // Explore edges from  $v$ .
7.   for each edge  $e=(v, u)$  in  $\text{adjList}[v]$ 
8.      $w = \text{weight of edge } e=(v, u)$ ;
    // If it's shorter to get to  $u$  from  $s$  via  $v$ , update.
9.     if  $\text{priority}[u] > \text{priority}[v] + w$ 
10.       $\text{priorityQueue.decreaseKey}(u,$ 
11.         $\text{priority}[v]+w)$ 
12.       $\text{predecessor}[u] = v$ 
13.    endif
14.  endwhile

```

15. Build SPT;

16. return SPT

Simple Examples




Works!

What if negative edges lengths?

-- What if run Dijkstra's on this?

-- Add constant to each edge?



Dijkstra's algorithm – proof of correctness

- Theorem: For every directed graph with non-negative edge lengths, Dijkstra's algorithm correctly computes all shortest path distances

Proof by induction:

Picture

Dijkstra's algorithm – proof of correctness

Proof by induction: mark nodes with a * once shortest path has been found

base case: - empty path is optimal

Inductive case: all previous iterations correct \rightarrow next iteration correct

1. Dijkstra picks next node to add to the shortest path tree w by finding the w with the smallest length path given by $\text{dist}(v^*) + l(v^*, w)$ where it has already found the shortest path to v^* .
---- must show this is the shortest path of all possible paths from s to w^* .
Know the shortest path to v^* has been found in a previous iteration by the inductive hypothesis
2. Suppose there is a shorter path, P_2 , s to w that has not been found yet. Any path from s to w must cross the frontier between nodes whose shortest paths are known and other nodes. Let the edge where P_2 crosses the frontier be (y^*, z) thus there is a path from s to y^* followed by the edge (y^*, z) , then a path from z to w .
3. But the the last section of P_2 must be ≥ 0 (since all edge lengths are non-negative).
4. Finally consider the path length of the part of P_2 from s to z given by $\text{dist}(s \text{ to } y^*) + l(y^*, z)$. But the path from s to z that is part of the proposed better path from s to w must be \leq shortest path from s to w . Why?. This contradicts that the vertex chosen by Dijkstra is the closest vertex to s among all the vertices one edge away from the vertices for which the shortest path is already known.

Copyright ©The McGraw-Hill Companies, Inc. Permission required for reproduction or display.
Kenneth H. Rosen, *Discrete Mathematics and its Applications*, 7e




FIGURE 4 Using Dijkstra's Algorithm to Find a Shortest Path from a to z .

Notes on Dijkstra's algorithm

- Doesn't work for graphs with negative weight
- Applicable to both undirected and directed graphs
- Efficiency
 - $O(|V|^2)$ for graphs represented by weight matrix and array implementation of priority queue
 - $O(|E|\log|V|)$ for graphs represented by adj. lists and min-heap implementation of priority queue
- Don't mix up Dijkstra's algorithm with Prim's algorithm!

Coding Problem

- Coding: assignment of bit strings to alphabet characters
- Codeword: bit string assigned to character
- Two types of codes:
 - fixed-length encoding (e.g., ASCII)
 - variable-length encoding (e.g., Morse code)
- Prefix-free codes: no codeword is a prefix of another codeword

Problem: If frequencies of the character occurrences are known, what is the best binary prefix-free code?

Morse and ASCII codes

Appendix B: ASCII Codes

Printable 8-bit ASCII codes

Decimal	Binary	Symbol	Decimal	Binary	Symbol	Decimal	Binary	Symbol
032	00100000	@	054	00100010	^	076	00100100	~
033	00100001	!	055	00100011	A	077	00100101	~
034	00100010	^	056	00100100	B	078	00100110	~
035	00100011	#	057	00100101	C	079	00100111	~
036	00100100	^	058	00100110	D	080	00100111	~
037	00100101	~	059	00100111	E	081	00101000	~
038	00100110	~	060	00101000	F	082	00101001	~
039	00100111	~	061	00101001	G	083	00101010	~
040	00101000	~	062	00101010	H	084	00101011	~
041	00101001	~	063	00101011	I	085	00101100	~
042	00101010	~	064	00101100	J	086	00101101	~
043	00101011	~	065	00101101	K	087	00101110	~
044	00101100	~	066	00101110	L	088	00101111	~
045	00101101	~	067	00101111	M	089	00110000	~
046	00101110	~	068	00101111	N	090	00110001	~
047	00101111	~	069	00110000	O	091	00110001	~
048	00110000	~	070	00110001	P	092	00110010	~
049	00110001	~	071	00110010	Q	093	00110011	~
050	00110010	~	072	00110011	R	094	00110100	~
051	00110011	~	073	00110100	S	095	00110101	~
052	00110100	~	074	00110101	T	096	00110110	~
053	00110101	~	075	00110110	U	097	00110111	~
054	00110110	~	076	00110111	V	098	00111000	~
055	00110111	~	077	00111000	W	099	00111001	~
056	00111000	~	078	00111001	X	100	00111010	~
057	00111001	~	079	00111010	Y	101	00111011	~
058	00111010	~	080	00111011	Z	102	00111100	~
059	00111011	~	081	00111100	~	103	00111101	~
060	00111100	~	082	00111101	~	104	00111110	~
061	00111101	~	083	00111110	~	105	00111111	~
062	00111110	~	084	00111111	~	106	00111111	~
063	00111111	~	085	00111111	~	107	00111111	~

191

University Publishing Online, hosted by Cambridge University Press © 2011

A •—
 B —•••
 C —•—
 D —••
 E •—
 F —•••
 G —•—
 H —•••
 I ••—

J •—•—
 K —•—
 L —•—•
 M —•—
 N —•—
 O —•—
 P —•—•
 Q —•—•
 R —•—•

S •••—
 T —•—
 U —•—
 V —•—•
 W —•—•
 X —•—•
 Y —•—•
 Z —•—•

Example

- Let $\Sigma = \{ \text{lower case letters, five punctuation marks and space} \}$
- How can we encode this – 5 bits since $2^5 = 32$
- Is there a way to reduce the length not of the code for each symbol BUT the average length of a message.
- Use frequency of the occurrence of the symbols
 - e, t, a - most frequent -- q, j, x, z – least frequent
 - Normalize so the sum of frequencies is = 1
 - Use frequencies to compute E(symbol length)
 - E.g. Morse code
- Prefix free codes are codes: for all symbols x and y – codeword(x) is not a prefix of codeword(y)
- Decoding is easy - - why?

Huffman codes – key insights

- Any binary tree with edges labeled with 0's and 1's yields a prefix-free code of characters assigned to its leaves
- Optimal binary tree minimizing the expected (weighted average) length of a codeword can be constructed as follows

.

Example

character	A	B	C	D	E
frequency	0.35	0.1	0.2	0.2	0.15

codeword	11	100	00	01	101
----------	----	-----	----	----	-----

average bits per character: 2.25

for fixed-length encoding: 3





compression ratio: $(3 - 2.25) / 3 * 100\% = 25\%$

Huffman's algorithm

- Initialize n one-node trees with alphabet characters and the tree weights with their frequencies.
- Repeat the following step $n-1$ times:
 - join two binary trees with smallest weights into one (as left and right subtrees)
 - make the new tree weight equal the sum of the weights of the two subtrees.
- Mark edges leading to left and right subtrees with 0's and 1's, respectively

Constructing a Huffman code tree

(a) f:5 e:9 c:12 b:13 d:16 a:45



Example: Build tree (smallest to left = 0)

character	A	B	C	D	E
frequency	0.32	0.25	0.2	0.18	0.05

Codewords:

H.C. average bits per character:

Fixed-length bits per character:

compression ratio =

Decode: 011110111011011


Huffman Code Proof Optimality: Big Picture

- Show that optimal prefix free code corresponds to a full binary tree.
- Huffman's algorithm puts two lowest frequency symbols as siblings of a full binary tree
- There exists an optimal code where the two symbols are at the lowest level of the tree, call these y and z
- Combining the two sibling symbols, y and z into w
 $ABL(T') + f_y + f_z = ABL(T)$ for any full binary tree
- Proof by induction: base case – 2 symbols– only one tree
 inductive case: true for k-1 symbols \rightarrow true for k symbols
 - » Assume $ABL(T_H) > ABL(T_{Opt}) \rightarrow ABL(T'_H) > ABL(T'_{Opt})$ gives contradiction using the above identity

Extensions and issues

- Image compression
 - Fraction of a bit for a white pixel, higher for black pixel
 - Video/audio – only send changes
- Adaptive encoding
- Many schemes are more effective for particular applications

Supplemental



Tree vertices	Remaining vertices	Illustration
$a(-, -)$	$b(a, 3)$ $c(-, \infty)$ $d(-, \infty)$ $e(a, 6)$ $f(a, 5)$	
$b(a, 3)$	$c(b, 1)$ $d(-, \infty)$ $e(a, 6)$ $f(b, 4)$	
$c(b, 1)$	$d(c, 6)$ $e(a, 6)$ $f(b, 4)$	
$f(b, 4)$	$d(f, 5)$ $e(f, 2)$	
$e(f, 2)$	$d(f, 5)$	
$d(f, 5)$		

FIGURE 9.3 Application of Prim's algorithm. The parenthesized labels of a vertex in the middle column indicate the nearest tree vertex and edge weight; selected vertices and edges are shown in bold.

Developing and remembering Prim's Algorithm

Big idea: grow a tree adding one vertex (closest) at a time -- Refinements

1. How find next vertex? -- what needs to happen to set up for finding the next vertex?
2. What are the operations needed?
3. What are possible data structures?
4. What should control flow look like?

Developing and remembering Prim's Algorithm - 2

Try to make it concrete

- Keep track of vertices in a list and associate with each the "edge connecting it to the tree and the edge weight"
- Store the graph either as an adjacency matrix or an adjacency list
- Initialize the list of vertices, mark the start vertex with nil, 0 since it has no edge and no weight
- Traverse the adjacency list of the start vertex and update distances and edges
- Loop till all the vertices are connected - $n-1$
 - remove the nearest vertex and mark it in the tree
 - update the distances and edges of vertices in the adjacency list

Single Source Shortest Paths Problem Dijkstra's algorithm

- Doesn't BFS solve shortest paths?
Yes, but only if the edge costs are all 1
- Why not just add edges to get the weight? (Good Question, can you use what you already know.)

Dijkstra's algorithm pseudo-code

single source all shortest paths


- Doesn't BFS solve shortest paths?
Yes, but only if the edge costs are all 1
- Why not just add edges to get the weight? (Good Question, can you use what you already know.)
- But problem is you may increase the number of edges significantly

Dijkstra's Algorithm: Implementation

- For each unexplored node, explicitly maintain $\pi(v) = \min_{e=(u,v): u \in S} d(u) + \ell_e$.
 - Next node to explore = node with minimum $\pi(v)$.
 - When exploring v , for each incident edge $e = (v, w)$, update

$$\pi(w) = \min \{ \pi(w), \pi(v) + \ell_e \}.$$
- Efficient implementation. Maintain a priority queue of unexplored nodes, prioritized by $\pi(v)$.

PQ Operation	Dijkstra	Array	Binary heap	d-way Heap	Fib heap [†]
Insert	n	n	log n	d log _d n	1
ExtractMin	n	n	log n	d log _d n	log n
ChangeKey	m	1	log n	log _d n	1
IsEmpty	n	1	1	1	1
Total		n ²	m log n	m log _{m/n} n	m + n log n



Tree vertices	Remaining vertices	Illustration
a(-, 0)	b(a, 3) c(-, ∞) d(a, 7) e(-, ∞)	
b(a, 3)	c(b, 3 + 4) d(b, 3 + 2) e(-, ∞)	
d(b, 5)	c(b, 7) e(d, 5 + 4)	
c(b, 7)	e(d, 9)	
e(d, 9)		

The shortest paths (identified by following nonnumeric labels backward from a destination vertex in the left column to the source) and their lengths (given by numeric labels of the tree vertices) are as follows:

from a to b: a - b of length 3
 from a to d: a - b - d of length 5
 from a to c: a - b - c of length 7
 from a to e: a - b - d - e of length 9

FIGURE 9.11 Application of Dijkstra's algorithm. The next closest vertex is shown in bold.

Another greedy algorithm for MST: Kruskal's

- Sort the edges in non-decreasing order of lengths
- "Grow" tree one edge at a time to produce MST through a series of expanding forests F_1, F_2, \dots, F_{n-1}
- On each iteration, add the next edge on the sorted list unless this would create a cycle. (If it would, skip the edge.)




FIGURE 9.5 Application of Kruskal's algorithm. Selected edges are shown in bold.

Notes about Kruskal's algorithm

- Algorithm looks easier than Prim's but is harder to implement (checking for cycles!)
- Cycle checking: a cycle is created if and only if added edge connects vertices in the same connected component
- Union-find* algorithms – see section 9.2

Correctness and Analysis: Kruskal's Algorithm

- Correctness : apply the cut property at each stage of the algorithm
- Analysis: Assuming n vertices and m edges. Using Union, Find and MakeSet operations discussed in the text with an efficient implementation
 - the algorithms order of complexity is dominated by the sorting operation on the edges hence it is $O(|E| * (\log |E|))$

Bellman Ford Algorithm

```

Bellman-Ford (G, w, s) // w: weights
  Initialize_Single_Source(V)
  For i = 1 to |V|-1
    for each edge (u,v) in E
      Relax((u,v), w)
  for each edge (u,v) in E
    if dist(v) > dist(u) + w(u,v)
      return false //negative cycle detected
  
```

Correctness of Bellman Ford Algorithm

- May want to defer till get to dynamic programming.
Good lab to write out as a dynamic programming algorithm
- Proof of correctness is proof of optimal substructure

DAG Algorithm

- DAG implies no cycles – no need to worry about negative weight cycles
- If we topologically sort the vertices, then any path must be consistent with the topological sort, that is the u comes before v in the path $\leftrightarrow u$ comes before v in the topological sort
- By making one pass over the vertices in the topological sort and relaxing the values of the vertices in their adjacency list we are guaranteed to find the shortest paths.

DAG algorithm pseudo code and efficiency


```

DAG-Shortest-Paths (G,w,s)
  Topologically sort vertices in G
  Initialize-Single-Source (G,s)
  For each vertex u // in topological sorted order
    For vertex v in the adjacency list of u
      Relax(u,v,w)
  
```


- Running time is $\theta(|E|+|V|)$ since topological sorting is $\theta(|E|+|V|)$ and the Relax function is called $\theta(|E|+|V|)$ times

Shortest Path Problem

- Shortest path network.
 - Directed graph $G = (V, E)$.
 - Source s , destination t .
 - Length ℓ_e = length of edge e .
 - cost of path = sum of edge costs in path
- Shortest path problem: find shortest directed path from s to t .



Cost of path s-2-3-5-t
 $= 9 + 23 + 2 + 16$
 $= 48.$



A* (heuristic): Brief comparison

- A* uses finds a least-cost path from a given initial node to one goal node (out of one or more possible goals).
- It uses a distance-plus-cost heuristic function $f(x) = g(x) + h(x)$ to determine the order in which the search visits nodes in the tree.
 - $g(x)$ – path cost function which is the cost from the starting node to the current node
 - $h(x)$ an **admissible** "heuristic estimate" of the distance to the goal (usually denoted). It must not over-estimate the distance to the goal